

EECS 484 Project 1: Fakebook Database

Due **Thursday, Sept 22 at 11:55 PM EST**

[Introduction](#)

[Submissions](#)

[Honor Code](#)

[Part 1 : Creating an ER Diagram \(62 points\)](#)

[Users](#)

[Messages](#)

[Photos and Albums](#)

[Events](#)

[Note on ER Diagram Design](#)

[Part 2: Creating the Data Tables \(50 points\)](#)

[Desired Schema](#)

[Part 3: Populate your database \(50 points\)](#)

[Sequences](#)

[Friends Trigger](#)

[Part 4: Creating External Views \(50 points\)](#)

[The Public Dataset](#)

[Oracle and SQL*Plus](#)

[Helpful SQL*PLUS Commands](#)

[Resetting Password and Sessions](#)

[Submitting](#)

[Appendix](#)

[Working with CAEN](#)

[Circular Dependencies for Foreign Keys](#)

[Debugging and Dependencies between Tables](#)

[FAQ From Past Semesters](#)

[SQL*Plus Potholes](#)

Introduction

In Project 1, you will be designing a relational database to store information for the fictional social media platform Fakebook. We will provide you with a description of the data you will need to store, complete with fields and requirements. Armed with the design specification, you will create an ER Diagram as well as a series of SQL scripts to create, load, and drop objects.

Submissions

This project is to be done in teams of 2 students (recommended) or individually. Both members of each team will receive the same score - the highest score among all submissions; as such, it is not necessary for each team member to submit the assignment. A tool for finding teammates has been made available on [Piazza](#). To create a team on the [Autograder](#), follow these steps **before making your first submission**:

1. One team member clicks the “Send group invitation” button on the Project 1 page.
2. The other team member confirms the invitation on their Autograder assignment page.

Do not make any submissions before joining your team! Once you click on “I’m working alone”, the Autograder will not let you change team members. If you do need to make a correction, the teaching staff has the ability to modify teams.

Project 1 is due on **Thursday, September 22 at 11:55 PM EST**. If you do not turn in your project by that deadline, or if you are unhappy with your work, you may continue to submit up until **Monday, September 26 at 11:55 PM EST** (4 days after the regular deadline). Please refer to [EECS 484 F22 Course Policies](#) for more information on *Assignments and Partners* and *Late Days*.

Honor Code

By submitting this project, you are agreeing to abide by the Honor Code: I have neither given nor received unauthorized aid on this assignment, nor have I concealed any violations of the Honor Code. You may not share answers with other students actively enrolled in the course outside of your teammate, nor may you consult with students who took the course in previous semesters. You are, however, allowed to discuss general approaches and class concepts with other students, and you are also permitted (and encouraged!) to post questions on Piazza.

Part 1 : Creating an ER Diagram (62 points)

Your first task of Project 1 is to design an ER Diagram that reflects the business rules of the Fakebook platform as described by the company's CEO Clark Huckelburg. Fakebook has four major features: *Users*, *Messages*, *Photos and Albums*, and *Events*. Descriptions of these features are listed below, though specifics such as data types and nullability are explicitly omitted. You may find later sections of this spec and/or the public dataset helpful in determining these specifics – practicing such design decisions is valuable for your growth as an engineer. **Do not make any additional assumptions, even if they would be reasonable in the “real world.”**

Users

Fakebook's *Users* feature is its most robust feature currently available to the public. When a Fakebook user signs up for the platform, they are assigned a **unique user ID**. Their profile also consists of a first name, last name, day, month, year of birth, and a non-binary gender. Additionally, users **may (but are not required to)** list a hometown city and/or a current city on their profile, and these cities can be changed at any time, though they **can only have** 1 hometown and 1 current city at any given time. Each city has a unique city ID, a city name, a state name, and a country name. The combination of city name, state name and country name is unique (you may not need to reflect this property in your ER Diagram).

In addition to its users' personal information, Fakebook maintains educational history on each user, which consists of **programs** and **graduation year**. Besides a unique program ID, each program also has a trio of fields: the institution (e.g. “University of Michigan”), the concentration (e.g. “Computer Science”) and the degree (e.g. “B.S.”); **this trio must be unique for every such program**. Users may list any number of programs (including 0) in their educational history; and a program may or may not be listed in the educational history of any number of users. Fakebook allows different users to register for the same program with different graduation years; however, **a user cannot list the same program multiple times with different graduation years**.

The last piece of the *Users* feature is friends. Two different Fakebook users can become friends through the platform, but a **user cannot befriend him/herself** (you may not need to reflect this property in your ER Diagram). Fakebook tracks the members of a friendship as “Requester” and “Requestee” (recall the concept of different roles in the same entity from lecture). There is no limit to the number of friends a Fakebook user can have. Also, a Fakebook user can have zero friends :/

Messages

Fakebook allows messages to be sent between two users (including themselves). Each message is given a unique message ID. Fakebook records the message content and the message sent time. It also tracks the user who sends the message as “Sender” and the user who receives the message as

“Receiver”. A Fakebook user can send or receive 0 or more messages, but **a message can only be sent exactly once (i.e. it has exactly one sender and exactly one receiver)**. Group messages are currently not supported by Fakebook.

Photos and Albums

Like any good social media platform, Fakebook allows its users to post photos. Once uploaded, photos are placed into albums and **each photo must belong to exactly one album**. Each photo is given a unique photo ID, and the metadata for photos consists of the photo uploaded time, last modified time of the photo, the photo’s link, and a caption. Fakebook does not directly track the owner of a photo, but this information can be retrieved via the album in which the photo is contained.

Each Fakebook album has a unique album ID and is **owned by exactly one** Fakebook user. There is no limit to the number of albums a single user can own, and there is no limit to the number of photos that an album can contain. However, each album **must contain at least one photo**. Fakebook tracks metadata for albums: the album name, the time the album was created, the last modified time of the album, the album’s link, and a visibility level (e.g. ‘Myself’, ‘Everyone’). In addition, each album must **have exactly one cover photo**; however, that **photo does not have to be one of the photos in the album**. A single photo can be the cover photo of 0 or more albums.

In addition to creating albums and uploading photos to those albums, Fakebook users can be tagged in the photos. Fakebook tracks the tagged user (but not the user doing the tagging), the tagged time, and the x-coordinate and y-coordinate within the photo. A user can be tagged in any number of photos, but **cannot be tagged in the same photo more than once**. A single photo can contain tags of 0 or more users. The tagged time and x, y coordinates may be the same or different (e.g., two tags on one photo could share the same x, y coordinates).

Events

The final feature of Fakebook is *Events*. An event itself is uniquely identified by an event ID and also contains a name, a tagline, and a description. Each event **is created by a single Fakebook user** (the “creator”); a Fakebook user can create 0 or more events. Other metadata for an event includes the host (not a Fakebook user but a simple string), the street address, the event’s type and subtype, the start and end time. Each event **must be located in exactly one city**; each city may have 0 or more events being held in.

The creator of an event does not have to participate in the event, which means that Fakebook events can have an unlimited number (including 0) of users participating. Each participant in an event has a **confirmation status** (e.g. ‘Attending’, ‘Declines’). Users can participate in any number of events, but no user can participate in the same event **more than once**, even with a different confirmation status.

Note on ER Diagram Design

Creating ER Diagrams is not an exact science: for a given specification, there are often several valid ways to represent all the necessary information in an ER Diagram. When grading your ER Diagrams, we will look to make sure that all of the entities, attributes, relations, keys, key constraints, and participation constraints are accurately depicted even if your diagram does not exactly match our intended solution. Also, note that **there may be some constraints described above that are not possible to depict on an ER Diagram**. As such, it is perfectly acceptable to ignore these constraints for Part 1; you'll implement them later in Part 2 instead.

Part 2: Creating the Data Tables (50 points)

Your second task of Project 1 is to write SQL DDL statements to create/drop data tables that reflect the Fakebook specifications. You will need to write 2 SQL scripts for this part: `createTables.sql` (to create the data tables) and `dropTables.sql` (to drop/destroy the data tables). These scripts should also create and drop any **constraints, sequences, and/or triggers** you find are necessary to enforce the rules of the Fakebook specification.

Once you have written these two files, you should run them within SQL*PLUS on your CAEN Linux machine. You should be able to run the commands below several times sequentially without error. If you cannot do this (i.e. if SQL*PLUS reports errors), you are liable to fail tests on the Autograder. To access your Oracle account, please refer to the section [Oracle and SQL*Plus](#).

```
SQL> @createTables.sql
SQL> @dropTables.sql
```

Desired Schema

We will test that your `createTables.sql` script properly creates the necessary data tables with all of the correct constraints. We will attempt to insert both valid and invalid data into your tables with the expectation that the valid inserts will be accepted and the invalid inserts will be rejected. To facilitate this, your tables must conform exactly to the schema below, even if it doesn't exactly match the schema you would have created following from your ER Diagram. You are not allowed to add any additional tables or columns to the schema, and both the column names and data types **must** match exactly. Deviating from this schema will cause you to fail tests on the Autograder.

Additionally, since we will be loading data into your tables, **to score full points in Part 2, you will need to understand Part 3** (especially the section on [Sequences](#) and [Friends Trigger](#)). You may also find it helpful to read the section on [Circular Dependencies for Foreign Keys](#).

USERS

- USER_ID (INTEGER) [Required field]
- FIRST_NAME (VARCHAR2(100)) [Required field]
- LAST_NAME (VARCHAR2(100)) [Required field]
- YEAR_OF_BIRTH (INTEGER)
- MONTH_OF_BIRTH (INTEGER)
- DAY_OF_BIRTH (INTEGER)
- GENDER (VARCHAR2(100))

FRIENDS

- USER1_ID (INTEGER) [Required field]
- USER2_ID (INTEGER) [Required field]

Important Note: This table should not allow duplicate friendships, regardless of the order in which the two IDs are listed. This means that (1, 9) and (9, 1) should be considered the *same* entry in this table – attempting to insert one while the other is already in the table should result in the insertion being rejected. The means of implementing this constraint is given later in the spec (see [Friends Trigger](#)).

CITIES

- CITY_ID (INTEGER) [Required field]
- CITY_NAME (VARCHAR2(100)) [Required field]
- STATE_NAME (VARCHAR2(100)) [Required field]
- COUNTRY_NAME (VARCHAR2(100)) [Required field]

USER_CURRENT_CITIES

- USER_ID (INTEGER) [Required field]
- CURRENT_CITY_ID (INTEGER) [Required field]

USER_HOMETOWN_CITIES

- USER_ID (INTEGER) [Required field]
- HOMETOWN_CITY_ID (INTEGER) [Required field]

MESSAGES

- MESSAGE_ID (INTEGER) [Required field]
- SENDER_ID (INTEGER) [Required field]
- RECEIVER_ID (INTEGER) [Required field]
- MESSAGE_CONTENT (VARCHAR2(2000)) [Required field]
- SENT_TIME (TIMESTAMP) [Required field]

PROGRAMS

- PROGRAM_ID (INTEGER) [Required field]
- INSTITUTION (VARCHAR2(100)) [Required field]

- CONCENTRATION (VARCHAR2(100)) [Required field]
- DEGREE (VARCHAR2(100)) [Required field]

EDUCATION

- USER_ID (INTEGER) [Required field]
- PROGRAM_ID (INTEGER) [Required field]
- PROGRAM_YEAR (INTEGER) [Required field]

USER_EVENTS

- EVENT_ID (INTEGER) [Required field]
- EVENT_CREATOR_ID (INTEGER) [Required field]
- EVENT_NAME (VARCHAR2(100)) [Required field]
- EVENT_TAGLINE (VARCHAR2(100))
- EVENT_DESCRIPTION (VARCHAR2(100))
- EVENT_HOST (VARCHAR2(100))
- EVENT_TYPE (VARCHAR2(100))
- EVENT_SUBTYPE (VARCHAR2(100))
- EVENT_ADDRESS (VARCHAR2(2000))
- EVENT_CITY_ID (INTEGER) [Required field]
- EVENT_START_TIME (TIMESTAMP)
- EVENT_END_TIME (TIMESTAMP)

PARTICIPANTS

- EVENT_ID (INTEGER) [Required field]
- USER_ID (INTEGER) [Required field]
- CONFIRMATION (VARCHAR2(100)) [Required field]

Confirmation must be one of these options (case-sensitive): ATTENDING, UNSURE, DECLINES, or NOT_REPLIED

ALBUMS

- ALBUM_ID (INTEGER) [Required field]
- ALBUM_OWNER_ID (INTEGER) [Required field]
- ALBUM_NAME (VARCHAR2(100)) [Required field]
- ALBUM_CREATED_TIME (TIMESTAMP) [Required field]
- ALBUM_MODIFIED_TIME (TIMESTAMP)
- ALBUM_LINK (VARCHAR2(100)) [Required field]
- ALBUM_VISIBILITY (VARCHAR2(100)) [Required field]

Album_visibility must be one of these options (case-sensitive): EVERYONE, FRIENDS, FRIENDS_OF_FRIENDS, or MYSELF

- COVER_PHOTO_ID (INTEGER) [Required field]

PHOTOS

- PHOTO_ID (INTEGER) [Required field]

- ALBUM_ID (INTEGER) [Required field]
- PHOTO_CAPTION (VARCHAR2(2000))
- PHOTO_CREATED_TIME (TIMESTAMP) [Required field]
- PHOTO_MODIFIED_TIME (TIMESTAMP)
- PHOTO_LINK (VARCHAR2(2000)) [Required field]

TAGS

- TAG_PHOTO_ID (INTEGER) [Required field]
- TAG_SUBJECT_ID (INTEGER) [Required field]
- TAG_CREATED_TIME (TIMESTAMP) [Required field]
- TAG_X (INTEGER) [Required field]
- TAG_Y (INTEGER) [Required field]

Feel free to use this schema to better inform the design of your ER Diagram, but do not feel like your diagram must represent this specific schema as long as all of the necessary constraints and other information are shown.

Don't forget to include things like **primary keys, foreign keys, NOT NULL requirements**, and other constraints to your DDLs even though they are not reflected in the schema list above. We recommend using your ER Diagram to assist in this.

Important Note: Using very long constraint names can cause some of your Autograder test cases to fail. Keep your constraint names short or don't use the **CONSTRAINT** keyword at all unless necessary. For example, inside of your **CREATE TABLE** statements, instead of writing

```
CONSTRAINT a_very_long_constraint_name CHECK (Column_A = 'A')
```

you can write

```
CHECK (Column_A = 'A')
```

Part 3: Populate your database (50 points)

The third part of Project 1 is to load the data from the public dataset into the tables you just created. To do this, you will have to write SQL DML statements that **SELECT** the appropriate data from the public dataset and **INSERT** that data into your tables. The names of the public tables, their fields, and a few business rules (input constraints) are listed at [The Public Dataset Schema](#), and they might give you some insight into how to design your ER Diagram and your own data tables. The public dataset is quite poorly designed, so you should not copy the public schema verbatim for your ER Diagram or you will lose a significant number of points.

You should put all of your DML statements into a single file named `loadData.sql` that loads data from the public dataset (and not from a private copy of that dataset). You are free to copy the public dataset to your own SQL*PLUS account for development and testing, but your scripts will not have access to this account when the Autograder runs them for testing.

When loading data for Fakebook friends, you should only include one directional pair of users even though Fakebook friendship is reciprocal. This means that if the public dataset includes both (2, 7) and (7, 2), only one of them (it doesn't matter which one) should be loaded into your table. The [Friends Trigger](#) will ensure that the direction of friendship matches what is expected, but you are still expected to properly select exactly one copy out of the public dataset.

Sequences

As you're loading data into your tables from the public dataset, you might find that you need ID numbers for entities where such ID numbers don't exist in the public data. The way to do this is to use a *sequence*, which is an SQL construct for generating streams of numbers. To create a sequence and use it to populate IDs for a table, use the following syntax, replacing the bracketed sections with the names/fields specific to your use case:

```
CREATE SEQUENCE [ sequence_name ]
    START WITH 1
    INCREMENT BY 1;

CREATE TRIGGER [ trigger_name ]
    BEFORE INSERT ON [ table_name ]
    FOR EACH ROW
    BEGIN
        SELECT [sequence_name].NEXTVAL INTO :NEW.[id_field] FROM DUAL;
    END;
/
```

Don't forget the trailing backslash!

Friends Trigger

Triggers are an SQL construct that can be used to execute arbitrary code when certain events happen, such as inserts into a table or updates of the contents of a table. You have already seen one trigger above, which we used to populate the ID field of a table when data is inserted.

In this project, you will also have to use a trigger to help enforce the more complicated constraint of the FRIENDS table. This trigger makes sure that **any incoming pair of friend IDs is sorted, which preserves uniqueness**. If you're having any difficulty understanding what this is doing, come to office hours and the staff will be happy to explain it. Because triggers are beyond the scope of this course, we have provided you with the entirety of the trigger syntax here:

```

CREATE TRIGGER Order_Friend_Pairs
  BEFORE INSERT ON FRIENDS
  FOR EACH ROW
    DECLARE temp INTEGER;
    BEGIN
      IF :NEW.USER1_ID > :NEW.USER2_ID THEN
        temp := :NEW.USER2_ID;
        :NEW.USER2_ID := :NEW.USER1_ID;
        :NEW.USER1_ID := temp;
      END IF ;
    END;
/

```

Don't forget the trailing backslash!

These sequences and triggers help define your database schema, and thus, are SQL DDL commands. Think about which file they should be included in.

Part 4: Creating External Views (50 points)

The final part of Project 1 is to create a set of external views for displaying the data you have loaded into your data tables. The views you create must have the exact same schema as the public dataset. This means that the column names and data types must match exactly; this schema is covered later in the spec. You will need to write 2 SQL scripts for this part: `createViews.sql` (to create the views and load data into them) and `dropViews.sql` (to drop/destroy the views). You should have a total of 5 views named as follows:

- VIEW_USER_INFORMATION
- VIEW_ARE_FRIENDS
- VIEW_PHOTO_INFORMATION
- VIEW_EVENT_INFORMATION
- VIEW_TAG_INFORMATION

Any use of the keyword “project1” in code or comment in your `createViews.sql` will cause your submission to automatically fail on the Autograder. This is to prevent any potential cheating. Please be cautious of this as you develop your solutions.

Once you have written these two files, you should be able to run them using SQL*PLUS from the command line of your CAEN Linux machine. You should be able to run the commands below several times sequentially without error. If you cannot do this (i.e. if SQL*PLUS reports errors), you are liable to fail tests on the Autograder.

```
SQL> @createTables.sql
SQL> @loadData.sql
SQL> @createViews.sql
SQL> @dropViews.sql
SQL> @dropTables.sql
```

For each of the views other than `VIEW_ARE_FRIENDS`, your views should exactly match the corresponding table in the public dataset. To test this, you can run the following queries in SQLplus, changing the name of the views as necessary. The output of both queries should be no rows selected; anything else indicates an error in your views.

```
SQL> SELECT * FROM project1.PUBLIC_USER_INFORMATION
2 MINUS
3 SELECT * FROM VIEW_USER_INFORMATION;
```

```
SQL> SELECT * FROM VIEW_USER_INFORMATION
2 MINUS
3 SELECT * FROM project1.PUBLIC_USER_INFORMATION;
```

To test `VIEW_ARE_FRIENDS`, use the following test scripts instead. The outputs should again be no rows selected.

```
SQL> SELECT LEAST(USER1_ID, USER2_ID), GREATEST(USER1_ID, USER2_ID)
2 FROM project1.PUBLIC_ARE_FRIENDS
3 MINUS
4 SELECT LEAST(USER1_ID, USER2_ID), GREATEST(USER1_ID, USER2_ID)
5 FROM VIEW_ARE_FRIENDS;
```

```
SQL> SELECT LEAST(USER1_ID, USER2_ID), GREATEST(USER1_ID, USER2_ID)
2 FROM VIEW_ARE_FRIENDS
3 MINUS
4 SELECT LEAST(USER1_ID, USER2_ID), GREATEST(USER1_ID, USER2_ID)
5 FROM project1.PUBLIC_ARE_FRIENDS;
```

The Public Dataset

The public dataset is divided into five tables, each of which has a series of data fields. Those data fields may or may not have additional business rules (constraints) that define the allowable values. When referring to any of these tables in your SQL scripts, you will need to use the fully-qualified table name by prepending `project1.` (including the “.”) to the table name (as seen in the testing examples above).

Here is an overview of the public dataset. All table names and field names are case-insensitive:

PUBLIC_USER_INFORMATION

1. USER_ID

The unique Fakebook ID of a user

2. FIRST_NAME

The user's first name; this is a required field

3. LAST_NAME

The user's last name; this is a required field

4. YEAR_OF_BIRTH

The year in which the user was born; this is an optional field

5. MONTH_OF_BIRTH

The month (as an integer) in which the user was born; this is an optional field

6. DAY_OF_BIRTH

The day on which the user was born; this is an optional field

7. GENDER

The user's gender; this is an optional field

8. CURRENT_CITY

The user's current city; this is an optional field, but if it is provided, so too will
CURRENT_STATE and CURRENT_COUNTRY

9. CURRENT_STATE

The user's current state; this is an optional field, but if it is provided, so too will
CURRENT_CITY and CURRENT_COUNTRY

10. CURRENT_COUNTRY

The user's current country; this is an optional field, but if it is provided, so too will
CURRENT_CITY and CURRENT_STATE

11. HOMETOWN_CITY

The user's hometown city; this is an optional field, but if it is provided, so too will
HOMETOWN_STATE and HOMETOWN_COUNTRY

12. HOMETOWN_STATE

The user's hometown state; this is an optional field, but if it is provided, so too will
HOMETOWN_CITY and HOMETOWN_COUNTRY

13. HOMETOWN_COUNTRY

The user's hometown country; this is an optional field, but if it is provided, so too will
HOMETOWN_CITY and HOMETOWN_STATE

14. INSTITUTION_NAME

The name of a college, university, or school that the user attended; this is an optional field,
but if it is provided, so too will PROGRAM_YEAR, PROGRAM_CONCENTRATION, and
PROGRAM_DEGREE

15. PROGRAM_YEAR

The year in which the user graduated from some college, university, or school; this is an
optional field, but if it is provided, so too will INSTITUTION_NAME,
PROGRAM_CONCENTRATION, and PROGRAM_DEGREE

16. PROGRAM_CONCENTRATION

The field in which the user studied at some college, university, or school; this is an optional field, but if it is provided, so too will INSTITUTION_NAME, PROGRAM_YEAR, and PROGRAM_DEGREE

17. PROGRAM_DEGREE

The degree the user earned from some college, university, or school; this is an optional field, but if it is provided, so too will INSTITUTION_NAME, PROGRAM_YEAR, and PROGRAM_CONCENTRATION

PUBLIC_ARE_FRIENDS

1. USER1_ID

The ID of the first of two Fakebook users in a friendship

2. USER2_ID

The ID of the second of two Fakebook users in a friendship

PUBLIC_PHOTO_INFORMATION

1. ALBUM_ID

The unique Fakebook ID of an album

2. OWNER_ID

The Fakebook ID of the user who owns the album

3. COVER_PHOTO_ID

The Fakebook ID of the album's cover photo

4. ALBUM_NAME

The name of the album; this is a required field

5. ALBUM_CREATED_TIME

The time at which the album was created; this is a required field

6. ALBUM_MODIFIED_TIME

The time at which the album was last modified; this is an optional field

7. ALBUM_LINK

The Fakebook URL of the album; this is a required field

8. ALBUM_VISIBILITY

The visibility/privacy level for the album; this is a required field

9. PHOTO_ID

The unique Fakebook ID of a photo in the album

10. PHOTO_CAPTION

The caption associated with the photo; this is an optional field

11. PHOTO_CREATED_TIME

The time at which the photo was created; this is a required field

12. PHOTO_MODIFIED_TIME

The time at which the photo was last modified; this is an optional field

13. PHOTO_LINK

The Fakebook URL of the photo; this is a required field

PUBLIC_TAG_INFORMATION

1. PHOTO_ID

The ID of a Fakebook photo

2. TAG_SUBJECT_ID

The ID of the Fakebook user being tagged in the photo

3. TAG_CREATED_TIME

The time at which the tag was created; this is a required field

4. TAG_X_COORDINATE

The x-coordinate of the location at which the subject was tagged; this is a required field

5. TAG_Y_COORDINATE

The y-coordinate of the location at which the subject was tagged; this is a required field

PUBLIC_EVENT_INFORMATION

1. EVENT_ID

The unique Fakebook ID of an event

2. EVENT_CREATOR_ID

The Fakebook ID of the user who created the event

3. EVENT_NAME

The name of the event; this is a required field

4. EVENT_TAGLINE

The tagline of the event; this is an optional field

5. EVENT_DESCRIPTION

A description of the event; this is an optional field

6. EVENT_HOST

The host of the event; this is an optional field, but it does not need to identify a Fakebook user

7. EVENT_TYPE

One of a predefined set of event types; this is an optional field, but the Fakebook front-end takes care of ensuring that the value is actually one of that predefined set by using a dropdown menu.

8. EVENT_SUBTYPE

One of a predefined set of event subtypes based on the event's type; this is an optional field, but cannot be provided unless the TYPE field is provided. The Fakebook front-end takes care of ensuring that the value is actually one of that predefined set by using a dropdown menu.

9. EVENT_ADDRESS

The street address at which the event is to be held; this is an optional field

10. EVENT_CITY

The city in which the event is to be held; this is a required field

11. EVENT_STATE

- The state in which the event is to be held; this is a required field
12. EVENT_COUNTRY
- The country in which the event is to be held; this is a required field
13. EVENT_START_TIME
- The time at which the event starts; this is an optional field
14. EVENT_END_TIME
- The time at which the event ends; this is an optional field

There is no data for event participants or messages in the public dataset, so you do not need to load anything into your table(s) corresponding to this information. However, you should assume that MESSAGE_ID would have been provided by the public dataset and does not need to be created using sequences and triggers.

Again, when referring to any of these tables in your SQL scripts, you will need to use the fully-qualified table name by prepending project1. (including the ".") to the table name.

Oracle and SQL*Plus

To access the public dataset for this project and to test your SQL scripts, you will be using a command line interface (CLI) from Oracle called SQL*PLUS. A SQL*PLUS account has been set up for you by the staff, so you should be all set to begin working on the project. If you encounter the error ORA-01017: invalid username/password; logon denied, then you may not have an account. Please privately post on Piazza with your username to request access.

To access your SQL*PLUS account, you must be on a CAEN Linux machine (see [Working with CAEN](#)). In order to use SQL*PLUS, you will need to load the class module by running `module load eecs484` in the command line. We suggest that you add this line to your `~/.bash_profile` so that it automatically runs every time you log in to your CAEN account. If you forget to load the module, you will encounter the error ORA-12162: TNS:net service name is incorrectly specified.

To start SQL*PLUS, type `sqlplus` at the command line and press enter; if you wish to have full access to your query history (with up/down arrows), type `rlwrap sqlplus` instead (we recommend using SQL*PLUS this way). Your username is your University of Michigan username, and your password is `eecclass` (this is case-sensitive). The first time you log in, the system will prompt you to change your password, which we recommend you do. Only use alphabetic characters, numerals, the underscore, the dollar sign, and the pound sign in your SQL*PLUS password.

Never use quotation marks or the @ symbol in your SQL*PLUS password. If you do, it is likely that you will not be able to log into your account, and you will need to reset it. See [Resetting Password and Sessions](#) for more details.

Helpful SQL*PLUS Commands

Once in SQL*PLUS, you can execute arbitrary SQL commands. You will notice that the formatting of output from SQL*PLUS can be less than ideal. Here are some tricks to make output more readable and some SQL commands to access information that might be important. Anything shown below in brackets should be replaced by an actual value:

- To view all of your tables, run:
`SELECT table_name FROM user_tables;`
- To view all of your views, run:
`SELECT view_name FROM user_views;`
- To view all of your sequences, run:
`SELECT sequence_name FROM user_sequences;`
- To view all of your triggers, run:
`SELECT trigger_name FROM user_triggers;`
- To view the full schema of any table, including the tables of the public dataset, run:
`DESC [table name];`
- To truncate the text in a particular column to only show a certain number of characters, run:
`COLUMN [column name] FORMAT a[num chars];`
- To remove the formatting from a particular column, run:
`cl [column name];`
and to remove the formatting from all columns, run:
`CLEAR COLUMNS;`
- To change the number of characters displayed on a single line from the default of 100, run:
`SET LINE [num chars];`
Another command that helps with formatting is:
`set markup csv on;`
- To select on the first several rows from a table you can use the ROWNUM pseudovvariable, such as:
`SELECT * FROM [table name] WHERE ROWNUM < [num];`
- To change your SQL*PLUS password, run:
`PASSWORD`
- To load commands in SQL*PLUS from a file, say `createTables.sql`, run (the name of the file is relative to the current directory from which `sqlplus` was launched):
`START createTables.sql`
or, run
`@createTables.sql`
- To quit SQL*PLUS, press ctrl+D or run:
`QUIT`

Resetting Password and Sessions

We have made an Autograder project named 'SQL*Plus Reset' that would allow you to reset your SQL*Plus password or kill your SQL*Plus sessions.

To reset your password to eecsclass, submit any text file named password.txt to this project.

To kill all your SQL*Plus sessions and fix the ORA-00054: resource busy and acquire with NOWAIT specified or timeout expired. error, submit any text file named sessions.txt to this project.

If you still have issues accessing your SQL*PLUS account after trying the solutions above, please email us and we can take a look. Keep in mind that this may take several hours, during which you will be unable to use SQL*PLUS to work on the project.

Submitting

There are two deliverables for Project 1: a PDF of your ER Diagram and your 5 SQL scripts (createTables.sql, dropTables.sql, loadData.sql, createViews.sql, and dropViews.sql). Part 1 (ER Diagram) is worth 62 points and Parts 2-4 (SQL scripts) are worth 50 points each, for a total of 212 points (62 on Gradescope, 150 on the Autograder).

Your ER Diagram will be submitted to [Gradescope](#) for hand-grading. The self-enrollment code for Gradescope is V5DP33 if you do not have access.

The PDF of your ER Diagram can be named whatever you would like. Your diagram can either be fully computer-generated or a scan of something hand-drawn. You may submit any number of times before the deadline. We will grade your latest submission. One team member should submit on Gradescope, but ***make sure to submit as a team, specifying your partner on Gradescope at submission time. If you do not do this, we will not be able to assign points to your partner.***

Your SQL scripts will be submitted to the [Autograder](#) for automated testing. Reach out via Piazza if you do not have access to the Autograder. It is your responsibility to ensure you have access *well before* the deadline.

Each group will be allowed 3 submissions per day with feedback; any submissions made in excess of those 3 will be graded, but the results of those submissions will be hidden from the group. Your highest scoring submission will be used for grading, with ties favoring your latest submission.

Appendix

Working with CAEN

To connect to CAEN remotely and to use SQL*Plus, you will need to SSH into a CAEN Linux machine. You will also need a way to sync your files between your local machine (to submit to the Autograder) and a CAEN machine (for SQL*PLUS testing). Many of you have visited the tutorial at [EECS 280 Setup Tutorial](#). You are free to use whatever setup has worked in the past for you, but we've summarized some options below.

Option 1 (Terminal):

The first option is to do all your development on your local machine, but periodically copy your files to CAEN for testing through the terminal.

You can sync your files with Git or a tool like rsync with the command `rsync -rtv <local_folder> <username>@login.engin.umich.edu:<remote_folder>`

Now that your files are on CAEN, to SSH into it, run `ssh <username>@login.engin.umich.edu`.

Option 2 (Visual Studio Code):

The second option is to utilize Visual Studio Code's SSH extension to do your development on a CAEN machine and sync your files back to your local machine with Git. This might be a better option if you want to test your scripts frequently, but it can take longer to debug machine-specific issues.

First, set up your SSH config file. An example `~/ssh/config` file might look like this:

```
# CAEN
Host caen
  Hostname login.engin.umich.edu
  User <username>
  ControlPersist 1h
```

Next, in Visual Studio Code, download the *Remote - SSH* extension. Under the extension settings, check the option to *Lockfiles in Tmp*. Open Command Palette and select the option *Connect to Host...* and choose *caen*. After entering your password and Duo authorization, you should now be able to edit your files directly on CAEN with Visual Studio Code as a text editor. To back up and sync your files to your local machine, you can use Git. For SSH multiplexing and more in-depth instructions, see [Avoiding Repeated 2FA](#).

Tip:

1. If you were able to successfully connect to caen **once** but it doesn't work after a little while/few days, try killing the connection first at least once and wait for the log notification to confirm that it was killed. To do so – open Command Palette and select the option *Remote - SSH: Kill VS Code Server on Host*.

Windows Users

For Windows users, we recommend using WSL as shown on the [EECS 280 website](#) and following the same steps as Linux/macOS users. You can, however, use tools like PuTTY, WinSCP, or Cyberduck to SSH into CAEN and sync your files.

Ultimately, you are free to use any setup you want as long as your files are formatted correctly when you submit to the Autograder. **Remember that running your scripts on CAEN will produce more helpful error descriptions than the ones provided by the Autograder.**

Circular Dependencies for Foreign Keys

Consider the following situation: you have two data tables, **TableA** and **TableB**. **TableA** needs to have a foreign key constraint on a column of **TableB**, and **TableB** needs to have a foreign key constraint on a column of **TableA**. How would you implement this in SQL?

One tempting solution is to directly include the foreign key constraints in your `CREATE TABLE` statements, but this unfortunately does not work. To create a foreign key, the table being referenced must already exist -- no matter which order we attempt to write out `CREATE TABLE` statements, the first one is going to fail because the other table will not yet exist.

Instead, we can add foreign key constraints to a table *after* it and the table it references have been created using an `ALTER TABLE` statement. The syntax for adding a foreign-key constraint to a previously created table is:

```
ALTER TABLE [ table_name ]  
ADD CONSTRAINT [ constraint_name ]  
[ constraint_syntax ]  
INITIALLY DEFERRED DEFERRABLE;
```

where “constraint syntax” should be the foreign key syntax that you would have put in a `CREATE TABLE` statement. (Note: the above `ALTER TABLE` syntax works for other kinds of SQL constraints as well.)

For simplicity and safety, you can write an `ALTER TABLE` statement using the above syntax for *both* tables in a circular dependency. You don't *have* to implement constraints this way to get full credit; other syntax can work and can be more concise. Remember that like sequences and triggers, `ALTER`

TABLE statements help define your database schema, and thus, are SQL DDL commands. Think about which file they should be included in.

Adding `INITIALLY DEFERRED DEFERRABLE` tells Oracle to defer the constraint check until later when you run your `loadData.sql` script to populate your tables. Why would we need to defer the check? Imagine your **Table A** and **Table B** have a circular dependency and are currently empty, but we are about to insert data into both tables. As soon as we insert data into **Table A** that is supposed to reference rows in **Table B**, Oracle will claim that **Table A**'s foreign key constraint has been violated, since all references to **Table B** in **Table A** don't currently point to anything valid (remember, **Table B** is currently empty!).

By *deferring* the check on **Table A**'s foreign key references to **Table B**, we can give Oracle the chance to insert data into both **Table A** and **Table B** before it performs the foreign key check. You can ensure that this happens *by grouping the INSERT statements for Table A and Table B into a single transaction*. A transaction is an atomic unit of work that the database performs. When you defer a constraint check, Oracle will wait until the end of a transaction to check for constraint violations. We will learn more about transactions at the end of the semester. By default, Oracle treats each individual SQL statement as its own transaction. This feature is called autocommitting, which is not desirable if you would like to insert into **Table A** and **Table B** in the same transaction.

Instead, you should manually define a transaction by turning off autocommit with the statement `SET AUTOCOMMIT OFF` in your script. After this statement, you can include your `INSERT` statements for **Table A** and **Table B**. After your `INSERT` statements, you should include the statement `COMMIT`. You will want to turn autocommit back on once this is done by including the statement `SET AUTOCOMMIT ON`.

Debugging and Dependencies between Tables

Be mindful of the dependencies between tables when debugging your code. For example, in Part 2, if your Users table fails to be created, other subsequent tables that depend on the Users table will also fail to be created. As another example, if you are failing `Test_User_Current_Cities` in Part 3 on the Autograder, check that you have first gotten `Test_Cities` correct.

FAQ From Past Semesters

Q: The order of columns in my table and/or view schemas does not match the order of columns in the public dataset's schema. Is this a problem?

A: No, this is not a problem. As long as the table names, column names, and column data types match, your schema will be valid.

Q: Are the IDs in the public dataset all unique?

A: Kind of. Each user/event/etc. in the public dataset has a unique ID, but there may be multiple

rows in a given table representing data for a single user/event/etc. In those cases, the IDs will be repeated.

Q: Do I need to include checking for the Type and Subtype fields in the Events table?

A: Nope.

Q: Can we trust all of the data in the public dataset?

A: All of the data in the public dataset conforms to all of the constraints laid out in this document. The only exception is the PUBLIC_ARE_FRIENDS table, which may contain impermissible duplicates.

Q: I looked up the schema for one of the tables, and I saw NUMBER(38) where the spec says the datatype should be INTEGER. Which should I use?

A: Our database uses INTEGER as an alias for a specifically-sized NUMBER type, which is why you see NUMBER(38) in the DESC output. Stick to using INTEGER in your DDLs.

Q: Is there an automatically-incrementing numeric type that I can use?

A: No, there is not. For those of you familiar with MySQL, Oracle has no equivalent to the auto increment specifier. You will have to use [Sequences](#) to achieve an equivalent effect.

Q: How do I make sure that every Album contains at least one Photo in my SQL scripts?

A: You can do this with a couple of more complicated triggers, but that is beyond the scope of this course, so you do not need to have this constraint enforced by your SQL scripts. You do, however, have to show this constraint on your ER diagram.

Q: There are many ways to declare a foreign key, how should I do it/is there an advised way?

A: For this project, it is advised that you use

```
FOREIGN KEY (Column_ID) REFERENCES Table_name (Other_ID)
```

Students have encountered autograder errors when attempting other declarations. If you have encountered errors while using a different declaration, it may be helpful to try this!

Q: Why do we sometimes use the AS keyword for aliases, but sometimes we don't?

A: In Oracle, column aliases can be defined as `column_name AS column_alias` or `column_name column_alias`. Table aliases can only be defined as `table_name table_alias`.

SQL*Plus Potholes

- Blank lines inside of a command will cancel the command.

This works

```
SELECT *  
FROM table_name;
```

But this will break your code

```
SELECT *  
  
FROM table_name;
```

- Begin multiline comment symbol `/*` must have a space right afterwards.

This works

```
SELECT * FROM table_name;  
/* this is a multiline  
comment block*/
```

But this will break your code

```
SELECT * FROM table_name;  
/*this is a multiline  
comment block*/
```

- Comments should not come right after the semicolon.

These work

```
SELECT * FROM table_name;  
-- this is a single line comment
```

```
SELECT * FROM table_name;  
/* this is a single line comment*/
```

But these will break your code

```
SELECT * FROM table_name; -- this is a single line comment
```

```
SELECT * FROM table_name; /* this is a single line comment*/
```